

[Dashboard](#) / [My courses](#) / [Computer Engineering & IT](#) / [CEIT-Even-sem-20-21](#) / [OS-Even-sem-2020-21](#) / [16 May - 22 May](#) / [End Sem Exam OS-2021](#)

Started on Saturday, 22 May 2021, 8:00 AM

State Finished

Completed on Saturday, 22 May 2021, 9:30 AM

Time taken 1 hour 30 mins

Grade 26.12 out of 40.00 (65%)

Question 1

Incorrect

Mark 0.00 out of 1.00

A 4 GB disk with 1 KB of block size would require these many number of **blocks** for it's free block bitmap:

Answer: ✖

The correct answer is: 512

Question 2

Correct

Mark 1.00 out of 1.00

Given that the memory access time is 110 ns, probability of a page fault is 0.5 and page fault handling time is 12 ms,
The effective memory access time in nanoseconds is:

Answer: ✔

The correct answer is: 6000055.00

Question 3

Incorrect

Mark 0.00 out of 1.00

The maximum size of a file in number of blocks of BSIZE in xv6 code is
(write a number only)

Answer: ✖

The correct answer is: 138

Question 4

Incorrect

Mark 0.00 out of 1.00

Calculate the average waiting time using
Round Robin scheduling with time quantum of 5 time units
for the following workload

assuming that they arrive in the order written below.

Process Burst Time

P1 5

P2 7

P3 6

P4 2

Write only a number in the answer upto two decimal points.

Answer: ✖

The correct answer is: 10.25



Question 5

Correct

Mark 1.00 out of 1.00

For the reference string

4 2 5 1 0 1 2 5 4 1 2

the number of page faults, including initial ones,
with FIFO replacement and 2 frames are :

Answer: ✓

4 -

4 2

5 2

5 1

0 1

-

2 1

2 5

4 5

4 1

2 1

The correct answer is: 10

Question 6

Correct

Mark 1.00 out of 1.00

Assuming a 16- KB page size, what is the page number for the address 428517 reference in decimal :
(give answer also in decimal)

Answer: ✓

The correct answer is: 26

Question 7

Correct

Mark 1.00 out of 1.00

In the code below assume that each function can be executed concurrently by many threads/processes.
Ignore syntactical issues, and focus on the semantics.

This program is an example of

```
spinlock a, b; // assume initialized
thread1() {
    spinlock(b);
    //some code;
    spinlock(a);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
thread2() {
    spinlock(a);
    //some code;
    spinlock(b);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
```

- ☒ a. Deadlock
- ☐ b. Self Deadlock
- ☐ c. None of these
- ☐ d. Deadlock or livelock depending on actual race
- ☐ e. Livelock



Your answer is correct.

The correct answer is: Deadlock

Question 8

Partially correct

Mark 1.33 out of 2.00

Match the snippets of xv6 code with the core functionality they achieve, or problems they avoid.
 "... means some code.

```
static inline uint
xchg(volatile uint *addr, uint newval)
{
    uint result;
```

// The + in "+m" denotes a read-modify-write operand.

```
asm volatile("lock; xchgl %0, %1" :
    "+m" (*addr), "=a" (result) :
    "1" (newval) :
    "cc");
return result;
}
```

Atomic compare and swap instruction (to be expanded inline into code)



```
void
sleep(void *chan, struct spinlock *lk)
{
    ...
    if(lk != &ptable.lock){
        acquire(&ptable.lock);
        release(lk);
    }
```

If you don't do this, a process may be running on two processors parallely



```
void
acquire(struct spinlock *lk)
{
    ...
    __sync_synchronize();
```

Tell compiler not to reorder memory access beyond this line



Your answer is partially correct.

You have correctly selected 2.

The correct answer is: static inline uint

```
xchg(volatile uint *addr, uint newval)
{
    uint result;
```

// The + in "+m" denotes a read-modify-write operand.

```
asm volatile("lock; xchgl %0, %1" :
    "+m" (*addr), "=a" (result) :
    "1" (newval) :
    "cc");
return result;
```

} → Atomic compare and swap instruction (to be expanded inline into code), void

```
sleep(void *chan, struct spinlock *lk)
```

```
{
```

```
...
```

```
if(lk != &ptable.lock){
    acquire(&ptable.lock);
    release(lk);
```

} → Avoid a self-deadlock, void

```
acquire(struct spinlock *lk)
```

```
{
```

```
...
```

```
__sync_synchronize(); → Tell compiler not to reorder memory access beyond this line
```

Question 9

Correct

Mark 1.00 out of 1.00

Predict the output of the program given here.

Assume that all the path names for the programs are correct. For example "/usr/bin/echo" will actually run echo command.

Assume that there is no mixing of printf output on screen if two of them run concurrently.

In the answer replace a new line by a single space.

For example::

good

output

should be written as good output

--

```
main() {
    int i;
    i = fork();
    if(i == 0)
        execl("/usr/bin/echo", "/usr/bin/echo", "hi", 0);
    else
        wait(0);
    fork();
    execl("/usr/bin/echo", "/usr/bin/echo", "one", 0);
}
```

Answer:



The correct answer is: hi one one

Question 10

Partially correct

Mark 1.67 out of 2.00

Select all the blocks that may need to be written back to disk (if updated, of-course), as "Yes", when an operation of deleting a file is carried out on ext2 file system.

An option has to be correct entirely to be marked "Yes"

Superblock



One or multiple data blocks of the parent directory



One or more data bitmap blocks for the parent directory



Block bitmap(s) for all the blocks of the file



Possibly one block bitmap corresponding to the parent directory



Data blocks of the file



Your answer is partially correct.

only one data block of parent directory. multiple blocks not possible. an entry is always contained within one single block

You have correctly selected 5.

The correct answer is: Superblock → Yes, One or multiple data blocks of the parent directory → No, One or more data bitmap blocks for the parent directory → No, Block bitmap(s) for all the blocks of the file → Yes, Possibly one block bitmap corresponding to the parent directory → Yes, Data blocks of the file → No



Question 11

Correct

Mark 1.00 out of 1.00

Select all the correct statements about bootloader.

Every wrong selection will deduct marks proportional to $1/n$ where n is total wrong choices in the question.

You will get minimum a zero.

- ☒ a. Modern Bootloaders often allow configuring the way an OS boots
- ☒ b. Bootloaders allow selection of OS to boot from
- ☐ c. Bootloader must be one sector in length
- ☐ d. The bootloader loads the BIOS
- ☒ e. LILO is a bootloader



Your answer is correct.

The correct answers are: LILO is a bootloader, Modern Bootloaders often allow configuring the way an OS boots, Bootloaders allow selection of OS to boot from



Question 12




Incorrect

Mark 0.00 out of 1.00


For each of the resource allocation diagram shown,
infer whether the graph contains at least one deadlock or not.

Yes

No



: Yes



Question 13

Partially correct

Mark 0.71 out of 1.00

Mark the statements about device drivers by marking as True or False.

True	False		
<input type="radio"/>	<input checked="" type="radio"/>	It's possible that a particular hardware has multiple device drivers available for it.	✗
<input checked="" type="radio"/>	<input type="radio"/>	xv6 has device drivers for IDE disk and console.	✓
<input checked="" type="radio"/>	<input type="radio"/>	A disk driver converts OS's logical view of disk into physical locations on disk.	✓
<input checked="" type="radio"/>	<input type="radio"/>	A device driver code is specific to a hardware device	✓
<input checked="" type="radio"/>	<input type="radio"/>	All devices of the same type (e.g. 2 hard disks) can typically use the same device driver	✓
<input type="radio"/>	<input checked="" type="radio"/>	Writing a device driver mandatorily demands reading the technical documentation about the hardware.	✗
<input type="radio"/>	<input checked="" type="radio"/>	Device driver is an intermediary between the end-user and OS	✓

It's possible that a particular hardware has multiple device drivers available for it.: True

xv6 has device drivers for IDE disk and console.: True

A disk driver converts OS's logical view of disk into physical locations on disk.: True

A device driver code is specific to a hardware device: True

All devices of the same type (e.g. 2 hard disks) can typically use the same device driver: True

Writing a device driver mandatorily demands reading the technical documentation about the hardware.: True

Device driver is an intermediary between the end-user and OS: False

Question 14

Partially correct

Mark 0.33 out of 1.00

Consider this program.

Some statements are identified using the // comment at the end.

Assume that = is an atomic operation.

```
#include <stdio.h>
#include <pthread.h>
long c = 0, c1 = 0, c2 = 0, run = 1;
void *thread1(void *arg) {
    while(run == 1) { //E
        c = 10; //A
        c1 = c2 + 5; //B
    }
}
void *thread2(void *arg) {
    while(run == 1) { //F
        c = 20; //C
        c2 = c1 + 3; //D
    }
}
int main() {
    pthread_t th1, th2;
    pthread_create(&th1, NULL, thread1, NULL);
    pthread_create(&th2, NULL, thread2, NULL);
    sleep(2);
    run = 0;
    fprintf(stdout, "c = %ld c1+c2 = %ld c1 = %ld c2 = %ld \n", c, c1+c2, c1, c2);
    fflush(stdout);
}
```

Which statements are part of the critical Section?

Yes	No	
<input checked="" type="radio"/>	<input checked="" type="radio"/>	F
<input checked="" type="radio"/>	<input type="radio"/>	D
<input checked="" type="radio"/>	<input checked="" type="radio"/>	C
<input checked="" type="radio"/>	<input checked="" type="radio"/>	A
<input checked="" type="radio"/>	<input type="radio"/>	B
<input checked="" type="radio"/>	<input checked="" type="radio"/>	E

F: No

D: Yes

C: No

A: No

B: Yes

E: No

Question 15

Partially correct

Mark 1.43 out of 2.00

Mark statements as T/F

All statements are in the context of preventing deadlocks.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	A process holding one resources and waiting for just one more resource can also be involved in a deadlock.	✓
<input type="radio"/>	<input checked="" type="radio"/>	If a resource allocation graph contains a cycle then there is a guarantee of a deadlock	✗
<input type="radio"/>	<input checked="" type="radio"/>	The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order	✗
<input checked="" type="radio"/>	<input type="radio"/>	Circular wait is avoided by enforcing a lock ordering	✓
<input checked="" type="radio"/>	<input type="radio"/>	Hold and wait means a thread/process holding some locks and waiting for acquiring some.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Deadlock is possible if all the conditions are met at the same time: Mutual exclusion, hold and wait, no pre-emption, circular wait.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens	✓

A process holding one resources and waiting for just one more resource can also be involved in a deadlock.: True

If a resource allocation graph contains a cycle then there is a guarantee of a deadlock: False

The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order: False

Circular wait is avoided by enforcing a lock ordering: True

Hold and wait means a thread/process holding some locks and waiting for acquiring some.: True

Deadlock is possible if all the conditions are met at the same time: Mutual exclusion, hold and wait, no pre-emption, circular wait.: True

Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens: True

Question 16

Correct

Mark 1.00 out of 1.00

Match the left side use(or non-use) of a synchronization primitive with the best option on the right side.

This is the smallest primitive made available in software, using the hardware provided atomic instructions	spinlock	✓
This tool is useful for event-wait scenarios	semaphore	✓
This tool is more useful on multiprocessor systems	spinlock	✓
This tool is quite attractive in solving the main bounded buffer problem	semaphore	✓
This tool is very useful for waiting for 'something'	condition variables	✓

Your answer is correct.

The correct answer is: This is the smallest primitive made available in software, using the hardware provided atomic instructions → spinlock, This tool is useful for event-wait scenarios → semaphore, This tool is more useful on multiprocessor systems → spinlock, This tool is quite attractive in solving the main bounded buffer problem → semaphore, This tool is very useful for waiting for 'something' → condition variables

Question 17

Correct

Mark 1.00 out of 1.00

The permissions -rwx--x--x on a file mean

- ☒ a. The file can be read only by the owner
- ☐ b. 'cat' on the file by owner will not work
- ☐ c. 'cat' on the file by any user will work
- ☒ d. 'rm' on the file by any user will work
- ☒ e. The file can be executed by anyone
- ☒ f. The file can be written only by the owner

✓

✓

✓

✓

Your answer is correct.

The correct answers are: The file can be executed by anyone, The file can be read only by the owner, The file can be written only by the owner, 'rm' on the file by any user will work

Question 18

Incorrect

Mark 0.00 out of 1.00

Note: for this question you get full marks if you select all and only correct options, you get ZERO if at least one option is wrong or not selected.

Select all the correct statements about log structured file systems.

- ☒ a. a transaction is said to be committed when all operations are written to file system
- ☒ b. log may be kept on same block device or another block device
- ☐ c. file system recovery may end up losing data
- ☒ d. even if file systems followed immediate writes (i.e. non-delayed writes), it could still require recovery
- ☒ e. file system recovery recovers all the lost data

✗

✓

✓

✗

Your answer is incorrect.

The correct answers are: file system recovery may end up losing data, log may be kept on same block device or another block device, even if file systems followed immediate writes (i.e. non-delayed writes), it could still require recovery

Question 19

Incorrect

Mark 0.00 out of 1.00

Consider the structure of directory entry in ext2, as shown in this diagram.

	inode	rec_len	file_type	name_len	name
0	21	12	1	2	· \0 \0 \0
12	22	12	2	2	· · \0 \0
24	53	16	5	2	h o m e 1 \0 \0 \0
40	67	28	3	2	u s r \0
52	0	16	7	1	o l d f i 1 e \0
68	34	12	4	2	s b i n

Select the correct statements about the directory entry in ext2 file system.

The correct formula for rec_len is (when entries are continuously stored)

- ☐ a. $\text{rec_len} = \text{sizeof}(\text{inode entry}) + \text{sizeof}(\text{name len entry}) + \text{sizeof}(\text{file type entry}) + (\text{strlen}(\text{name}) + (-1) * (\text{strlen}(\text{name}) \% 4)$
- ☐ b. $\text{rec_len} = \text{sizeof}(\text{inode entry}) + \text{sizeof}(\text{name len entry}) + \text{sizeof}(\text{file type entry}) + (\text{strlen}(\text{name}) + (\text{strlen}(\text{name}) - 4) \% 4$
- ☒ c. $\text{rec_len} = \text{sizeof}(\text{inode entry}) + \text{sizeof}(\text{name len entry}) + \text{sizeof}(\text{file type entry}) + (\text{strlen}(\text{name}) + 4 - (\text{strlen}(\text{name}) \% 4)$
- ☐ d. $\text{rec_len} = \text{sizeof}(\text{inode entry}) + \text{sizeof}(\text{name len entry}) + \text{sizeof}(\text{file type entry}) + (\text{strlen}(\text{name}) + (-1) * (\text{strlen}(\text{name}) - 4)$
- ☐ e. $\text{rec_len} = \text{sizeof}(\text{inode entry}) + \text{sizeof}(\text{name len entry}) + \text{sizeof}(\text{file type entry}) + (\text{strlen}(\text{name}) \% 4$
- ☐ f. $\text{rec_len} = \text{sizeof}(\text{inode entry}) + \text{sizeof}(\text{name len entry}) + \text{sizeof}(\text{file type entry}) + \text{strlen}(\text{name})$

Your answer is incorrect.

The correct answer is: $\text{rec_len} = \text{sizeof}(\text{inode entry}) + \text{sizeof}(\text{name len entry}) + \text{sizeof}(\text{file type entry}) + (\text{strlen}(\text{name}) + (-1) * (\text{strlen}(\text{name}) - 4)$

Question 20

Partially correct

Mark 0.50 out of 1.00

Mark whether the given sequence of events is possible or not-possible. Also, select the reason for your answer.

For each sequence it's a not-possible sequence if some important event is not mentioned in the sequence.

Assume that the kernel code is non-interruptible and uniprocessor system.

Process P1 executing a system call

Timer interrupt

Generic interrupt handler runs

Scheduler runs

Scheduler selects P2 for execution

P2 returns from timer interrupt handler

Process p2, user code executing

This sequence of events is: ✓

Because

✗

Question 21

Incorrect

Mark 0.00 out of 1.00

The given semaphore implementation faces which problem?

Assume any suitable code for signal()

Note: blocks means waits in a wait queue.

```
struct semaphore {
    int val;
    spinlock lk;
};
sem_init(semaphore *s, int initval) {
    s->val = initval;
    s->sl = 0;
}
wait(semaphore *s) {
    spinlock(&(s->sl));
    while(s->val <= 0)
        ;
    (s->val)--;
    spinunlock(&(s->sl));
}
```

- ☐ a. blocks holding a spinlock
- ☐ b. deadlock
- ☒ c. too much spinning, bounded wait not guaranteed
- ☐ d. not holding lock after unblock

✖

Your answer is incorrect.

The correct answer is: deadlock

Question 22

Partially correct

Mark 0.80 out of 1.00

Mark statements True/False w.r.t. change of states of a process.

Reference: The process state diagram (and your understanding of how kernel code works)




Figure 3.2 Diagram of process state.

True	False		
<input type="radio"/>	<input checked="" type="radio"/>	A process in RUNNING state only can become TERMINATED because scheduler moves it to ZOMBIE state	✓
<input checked="" type="radio"/>	<input type="radio"/>	A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.	✗
<input checked="" type="radio"/>	<input type="radio"/>	A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred	✓
<input checked="" type="radio"/>	<input type="radio"/>	Every process has to go through ZOMBIE state, at least for a small duration.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Only a process in READY state is considered by scheduler	✓

A process in RUNNING state only can become TERMINATED because scheduler moves it to ZOMBIE state: False

A process in READY state can not go to WAITING state because the resource on which it will WAIT will not be in use when process is in READY state.: False

A process in WAITING state can not become RUNNING because the event it's waiting for has not occurred: True

Every process has to go through ZOMBIE state, at least for a small duration.: True

Only a process in READY state is considered by scheduler: True

Question 23

Correct

Mark 1.00 out of 1.00

Select T/F for statements about Volume Managers.

Do pay attention to the use of the words physical partition and physical volume.

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The volume manager can create further internal sub-divisions of a physical partition for efficiency or features.	✓
<input checked="" type="radio"/>	<input type="radio"/>	A logical volume can be extended in size but upto the size of volume group	✓
<input checked="" type="radio"/>	<input type="radio"/>	A logical volume may span across multiple physical volumes	✓
<input checked="" type="radio"/>	<input type="radio"/>	The volume manager stores additional metadata on the physical disk partitions	✓
<input checked="" type="radio"/>	<input type="radio"/>	A physical partition should be initialized as a physical volume, before it can be used by volume manager.	✓
<input checked="" type="radio"/>	<input type="radio"/>	A volume group consists of multiple physical volumes	✓
<input checked="" type="radio"/>	<input type="radio"/>	A logical volume may span across multiple physical partitions	✓ since a physical volume is made up of physical partitions, and a volume can span across multiple PVs, it can also span across multiple PP

The volume manager can create further internal sub-divisions of a physical partition for efficiency or features.: True

A logical volume can be extended in size but upto the size of volume group: True

A logical volume may span across multiple physical volumes: True

The volume manager stores additional metadata on the physical disk partitions: True

A physical partition should be initialized as a physical volume, before it can be used by volume manager.: True

A volume group consists of multiple physical volumes: True

A logical volume may span across multiple physical partitions: True

Question 24

Correct

Mark 1.00 out of 1.00

Map the block allocation scheme with the problem it suffers from

(Match pairs 1-1, match a scheme with the problem that it suffers from relatively the most, compared to others)

Continuous allocation	need for compaction	✓
Linked allocation	Too many seeks	✓
Indexed Allocation	Overhead of reading metadata blocks	✓

Your answer is correct.

The correct answer is: Continuous allocation → need for compaction, Linked allocation → Too many seeks, Indexed Allocation → Overhead of reading metadata blocks

Question 25

Correct

Mark 1.00 out of 1.00

This one is not a system call:

- ☐ a. open
☐ b. read
☐ c. write
☒ d. scheduler

✓

Your answer is correct.

The correct answer is: scheduler



Question 26

Correct

Mark 1.00 out of 1.00

Match the pairs.

This question is based on your general knowledge about operating systems/related concepts and their features.

Java threads	monitors, re-entrant locks, semaphores	✓
Linux threads	atomic-instructions, spinlocks, etc.	✓
POSIX threads	semaphore, mutex, condition variables	✓

Your answer is correct.

The correct answer is: Java threads → monitors, re-entrant locks, semaphores, Linux threads → atomic-instructions, spinlocks, etc., POSIX threads → semaphore, mutex, condition variables

Question 27

Correct

Mark 1.00 out of 1.00

Consider the following list of free chunks, in continuous memory management:

7k, 15k, 21k, 14k, 19k, 6k

Suppose there is a request for chunk of size 5k, then the free chunk selected under each of the following schemes will be

Best fit:	6k	✓
First fit:	7k	✓
Worst fit:	21k	✓

Question 28

Correct

Mark 1.00 out of 1.00

This one is not a scheduling algorithm

- ☐ a. Round Robin
- ☐ b. SJF
- ☒ c. Mergesort
- ☐ d. FCFS

✓

Your answer is correct.

The correct answer is: Mergesort

Question 29

Correct

Mark 1.00 out of 1.00

Mark whether the concept is related to scheduling or not.

Yes	No		
<input checked="" type="radio"/>	<input type="radio"/>	timer interrupt	✓
<input checked="" type="radio"/>	<input type="radio"/>	context-switch	✓
<input checked="" type="radio"/>	<input type="radio"/>	ready-queue	✓
<input type="radio"/>	<input checked="" type="radio"/>	file-table	✓
<input checked="" type="radio"/>	<input type="radio"/>	runnable process	✓

timer interrupt: Yes

context-switch: Yes

ready-queue: Yes

file-table: No

runnable process: Yes

Question 30

Partially correct

Mark 1.00 out of 2.00

Map ext2 data structure features with their purpose

Many
copies of
Superblock

Choose...

Free
blocks
count in
superblock
and group
descriptor

Redundancy to ensure the most crucial data structure is not lost

Used
directories
count in
group
descriptor

is redundant and helps do calculations of directory entries faster

Combining
file type
and
access
rights in
one
variable

saves 1 byte of space

rec_len
field in
directory
entry

Try to keep all the data of a directory and it's file close together in a group

File Name
is padded

aligns all memory accesses on word boundary, improving performance

Inode
bitmap is
one block

limits total number of files that can belong to a group

Block
bitmap is
one block

Limits the size of a block group, thus improvising on purpose of a group

Mount
count in
superblock

to enforce file check after certain amount of mounts at boot time

Inode table
location in
Group
Descriptor

is redundant and helps do calculations of directory entries faster

Inode table

All inodes are kept together so that one disk read leads to reading many inodes together, effectively doing a buffering of subsequent inode reads, and to save space on disk

A group

Redundancy to ensure the most crucial data structure is not lost

Your answer is partially correct.

You have correctly selected 6.

The correct answer is: **Many copies of Superblock** → Redundancy to ensure the most crucial data structure is not lost, **Free blocks count in superblock and group descriptor** → Redundancy to help fsck restore consistency, **Used directories count in group descriptor** → attempt is made to evenly spread the first-level directories, this count is used there, **Combining file type and access rights in one variable** → saves 1 byte of space, **rec_len field in directory entry** → allows holes and linking of entries in directory, File Name is padded → aligns all memory accesses on word boundary, improving performance, **Inode bitmap is one block** → limits total number of files that can belong to a group, **Block bitmap is one block** → Limits the size of a block group, thus improvising on purpose of a group, **Mount count in superblock** → to enforce file check after certain amount of mounts at boot time, **Inode table location in Group Descriptor** → Obvious, as it's per group and not per file-system, **Inode table** → All inodes are kept together so that one disk read leads to reading many inodes together, effectively doing a buffering of subsequent inode reads, and to save space on disk, **A group** → Try to keep all the data of a directory and it's file close together in a group



Question 31

Partially correct

Mark 1.85 out of 2.00

Mark True/False

Statements about scheduling and scheduling algorithms

True	False		
<input checked="" type="radio"/>	<input type="radio"/>	The nice() system call is used to set priorities for processes	✓
<input checked="" type="radio"/>	<input type="radio"/>	Aging is used to ensure that low-priority processes do not starve in priority scheduling.	✓
<input type="radio"/>	<input checked="" type="radio"/>	In non-pre-emptive priority scheduling, the highest priority process is scheduled and runs until it gives up CPU.	✗
<input checked="" type="radio"/>	<input type="radio"/>	xv6 code does not care about Processor Affinity	✓
<input checked="" type="radio"/>	<input type="radio"/>	In pre-emptive priority scheduling, priority is implemented by assigning more time quantum to higher priority process.	✓
<input checked="" type="radio"/>	<input type="radio"/>	A scheduling algorithm is non-preemptive if it does context switch only if a process voluntarily relinquishes CPU or it terminates.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Processor Affinity refers to memory accesses of a process being stored on cache of that processor	✓
<input checked="" type="radio"/>	<input type="radio"/>	Response time will be quite poor on non-interruptible kernels	✓
<input checked="" type="radio"/>	<input type="radio"/>	Shortest Remaining Time First algorithm is nothing but pre-emptive Shortest Job First algorithm	✓
<input checked="" type="radio"/>	<input type="radio"/>	On Linuxes the CPU utilisation is measured as the time spent in scheduling the idle thread	✓
<input checked="" type="radio"/>	<input type="radio"/>	Generally the voluntary context switches are much more than non-voluntary context switches on a Linux system.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Pre-emptive scheduling leads to many race conditions in kernel code.	✓
<input checked="" type="radio"/>	<input type="radio"/>	Statistical observations tell us that most processes have large number of small CPU bursts and relatively smaller numbers of large CPU bursts.	✓

The nice() system call is used to set priorities for processes: True

Aging is used to ensure that low-priority processes do not starve in priority scheduling.: True

In non-pre-emptive priority scheduling, the highest priority process is scheduled and runs until it gives up CPU.: True

xv6 code does not care about Processor Affinity: True

In pre-emptive priority scheduling, priority is implemented by assigning more time quantum to higher priority process.: True

A scheduling algorithm is non-preemptive if it does context switch only if a process voluntarily relinquishes CPU or it terminates.: True

Processor Affinity refers to memory accesses of a process being stored on cache of that processor: True

Response time will be quite poor on non-interruptible kernels: True

Shortest Remaining Time First algorithm is nothing but pre-emptive Shortest Job First algorithm: True

On Linuxes the CPU utilisation is measured as the time spent in scheduling the idle thread: True

Generally the voluntary context switches are much more than non-voluntary context switches on a Linux system.: True

Pre-emptive scheduling leads to many race conditions in kernel code.: True

Statistical observations tell us that most processes have large number of small CPU bursts and relatively smaller numbers of large CPU bursts.: True



Question 32

Partially correct

Mark 1.17 out of 2.00

The unix file semantics demand that changes to any open file are visible immediately to any other processes accessing that file at that point in time.

Select the data-structure/programmatic features that ensure the implementation of unix semantics. (Assume there is no mmap())

Yes	No		
<input type="radio"/>	<input checked="" type="radio"/>	All processes accessing the same file share the file descriptor among themselves	✓
<input type="radio"/>	<input checked="" type="radio"/>	The pointer entry in the file descriptor array entry points to the data of the file directly	✓
<input checked="" type="radio"/>	<input type="radio"/>	There is only one global file structure per on-disk file.	✗
<input type="radio"/>	<input checked="" type="radio"/>	All file accesses are made using only global variables	✓
<input checked="" type="radio"/>	<input type="radio"/>	The 'file offset' is shared among all the processes that access the file.	✗
<input type="radio"/>	<input checked="" type="radio"/>	No synchronization is implemented so that changes are made available immediately.	✓
<input checked="" type="radio"/>	<input type="radio"/>	A single spinlock is to be used to protect the unique global 'file structure' representing the file, thus synchronizing access, and making other processes wait for earlier process to finish writing so that writes get visible immediately.	✗
<input checked="" type="radio"/>	<input type="radio"/>	There is only one in-memory copy of the on disk file's contents in kernel memory/buffers	✓
<input checked="" type="radio"/>	<input type="radio"/>	The file descriptors in every PCB are pointers to the same global file structure.	✗
<input type="radio"/>	<input checked="" type="radio"/>	The file descriptor array is external to PCB and all processes that share a file, have pointers to same file-descriptors' array	✓
<input checked="" type="radio"/>	<input type="radio"/>	All file structures representing any open file, give access to the same in-memory copy of the file's contents	✓
<input checked="" type="radio"/>	<input type="radio"/>	The 'file offset' index is stored outside the file-structure to which file-descriptor array points	✗

All processes accessing the same file share the file descriptor among themselves: No

The pointer entry in the file descriptor array entry points to the data of the file directly: No

There is only one global file structure per on-disk file.: No

All file accesses are made using only global variables: No

The 'file offset' is shared among all the processes that access the file.: No

No synchronization is implemented so that changes are made available immediately.: No

A single spinlock is to be used to protect the unique global 'file structure' representing the file, thus synchronizing access, and making other processes wait for earlier process to finish writing so that writes get visible immediately.: No

There is only one in-memory copy of the on disk file's contents in kernel memory/buffers: Yes

The file descriptors in every PCB are pointers to the same global file structure.: No

The file descriptor array is external to PCB and all processes that share a file, have pointers to same file-descriptors' array: No

All file structures representing any open file, give access to the same in-memory copy of the file's contents: Yes

The 'file offset' index is stored outside the file-structure to which file-descriptor array points: No

Question 33

Partially correct

Mark 0.33 out of 2.00

Map the function in xv6's file system code, to it's perceived logical layer.

namei	inode	✗
filestat()	Choose...	
dirlookup	directory	✓
ialloc	file descriptor	✗
stati	Choose...	
ideintr	buffer cache	✗
bread	Choose...	
balloc	file descriptor	✗
sys_chdir()	system call	✓
skipelem	system call	✗
commit	system call	✗
bmap	system call	✗

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: namei → pathname lookup, filestat() → file descriptor, dirlookup → directory, ialloc → inode, stati → inode, ideintr → disk driver, bread → buffer cache, balloc → block allocation on disk, sys_chdir() → system call, skipelem → pathname lookup, commit → logging, bmap → inode

[◀ Course Exit Feedback](#)

Jump to...

[xv6-public-master ▶](#)