

| | |
|---------------------|---|
| Started on | Saturday, 16 March 2024, 1:32 PM |
| State | Finished |
| Completed on | Saturday, 16 March 2024, 3:33 PM |
| Time taken | 2 hours 1 min |
| Grade | 13.35 out of 15.00 (88.97%) |

Question **1**

Correct

Mark 0.50 out of 0.50

Select the correct statements about sched() and scheduler() in xv6 code

- ☒ a. Each call to sched() or scheduler() involves change of one stack inside swtch() ✓
- ☒ b. When either sched() or scheduler() is called, it does not return immediately to caller ✓
- ☒ c. sched() switches to the scheduler's context ✓
- ☒ d. sched() and scheduler() are co-routines ✓
- ☒ e. When either sched() or scheduler() is called, it results in a context switch ✓
- ☒ f. After call to swtch() in scheduler(), the control moves to code in sched() ✓
- ☒ g. scheduler() switches to the selected process's context ✓
- ☒ h. After call to swtch() in sched(), the control moves to code in scheduler() ✓

Your answer is correct.

The correct answers are: sched() and scheduler() are co-routines, When either sched() or scheduler() is called, it does not return immediately to caller, When either sched() or scheduler() is called, it results in a context switch, sched() switches to the scheduler's context, scheduler() switches to the selected process's context, After call to swtch() in scheduler(), the control moves to code in sched(), After call to swtch() in sched(), the control moves to code in scheduler(), Each call to sched() or scheduler() involves change of one stack inside swtch()

Question **2**

Correct

Mark 0.50 out of 0.50

The variable \$stack in entry.S is

- ☐ a. located at 0x7c00
- ☐ b. located at the value given by %esp as setup by bootmain()
- ☒ c. a memory region allocated as a part of entry.S ✓
- ☐ d. located at less than 0x7c00
- ☐ e. located at 0

The correct answer is: a memory region allocated as a part of entry.S

Question 3

Correct

Mark 0.50 out of 0.50

Consider the image given below, which explains how paging works.

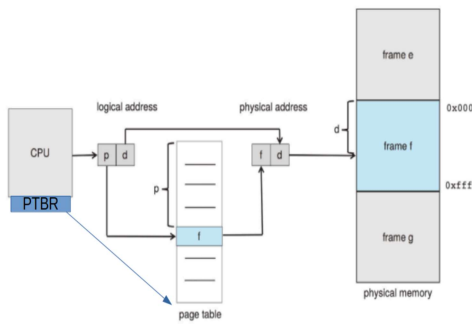


Figure 9.8 Paging hardware.

Mention whether each statement is True or False, with respect to this image.

| True | False | | |
|----------------------------------|----------------------------------|---|---|
| <input type="radio"/> | <input checked="" type="radio"/> | Size of page table is always determined by the size of RAM | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | The PTBR is present in the CPU as a register | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | Maximum Size of page table is determined by number of bits used for page number | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | The page table is indexed using frame number | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | The page table is itself present in Physical memory | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | The page table is indexed using page number | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | The physical address may not be of the same size (in bits) as the logical address | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | The locating of the page table using PTBR also involves paging translation | ✓ |

Size of page table is always determined by the size of RAM: False

The PTBR is present in the CPU as a register: True

Maximum Size of page table is determined by number of bits used for page number: True

The page table is indexed using frame number: False

The page table is itself present in Physical memory: True

The page table is indexed using page number: True

The physical address may not be of the same size (in bits) as the logical address: True

The locating of the page table using PTBR also involves paging translation: False

Question 4

Correct

Mark 0.50 out of 0.50

Mark statements as True/False w.r.t. the creation of free page list in xv6.

| True | False | |
|----------------------------------|----------------------------------|---|
| <input checked="" type="radio"/> | <input type="radio"/> | the kmem.lock is used by kfree() and kalloc() only. ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | free page list is a singly circular linked list. ✓ it's singly linked NULL terminated list. |
| <input checked="" type="radio"/> | <input type="radio"/> | if(kmem.use_lock) acquire(&kmem.lock); is not done when called from kinit1() because there is no need to take the lock when kinit1() is running because interrupts are disabled and only one processor is running ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | kmem.use_lock is set to 1 after free page list is created, so that kmem.lock is taken before accessing kmem.freelist. ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | if(kmem.use_lock) acquire(&kmem.lock); this "if" condition is true, when kinit2() runs because multi-processor support has been enabled by now. ✓ No. kinit2() calls kfree() and then initializes use_lock. |
| <input checked="" type="radio"/> | <input type="radio"/> | The pointers that link the pages together are in the first 4 bytes of the pages themselves ✓ |

the kmem.lock is used by kfree() and kalloc() only.: True

free page list is a singly circular linked list.: False

if(kmem.use_lock)

 acquire(&kmem.lock);

is not done when called from kinit1() because there is no need to take the lock when kinit1() is running because interrupts are disabled and only one processor is running: True

kmem.use_lock is set to 1 after free page list is created, so that kmem.lock is taken before accessing kmem.freelist.: True

if(kmem.use_lock)

 acquire(&kmem.lock);

this "if" condition is true, when kinit2() runs because multi-processor support has been enabled by now.: False

The pointers that link the pages together are in the first 4 bytes of the pages themselves: True

Question **5**

Correct

Mark 0.50 out of 0.50

Map ext2 data structure features with their purpose

| | |
|---|--|
| Mount count in superblock ✓ | to enforce file check after certain amount of mounts at boot time |
| A group ✓ | Try to keep all the data of a directory and it's file close together in a group |
| Combining file type and access rights in one variable ✓ | saves 1 byte of space |
| Block bitmap is one block ✓ | Limits the size of a block group, thus improvising on purpose of a group |
| Inode bitmap is one block ✓ | limits total number of files that can belong to a group |
| File Name is padded ✓ | aligns all memory accesses on word boundary, improving performance |
| rec_len field in directory entry ✓ | allows holes and linking of entries in directory |
| Inode table location in Group Descriptor ✓ | Obvious, as it's per group and not per file-system |
| Inode table ✓ | All inodes are kept together so that one disk read leads to reading many inodes together, effectively doing a buffering of subsequent inode reads, and to save |
| Used directories count in group descriptor ✓ | attempt is made to evenly spread the first-level directories, this count is used there |
| Free blocks count in superblock and group descriptor ✓ | Redundancy to help fsck restore consistency |
| Many copies of Superblock ✓ | Redundancy to ensure the most crucial data structure is not lost |

Your answer is correct.

The correct answer is: **Mount count in superblock** → to enforce file check after certain amount of mounts at boot time, **A group** → Try to keep all the data of a directory and it's file close together in a group, **Combining file type and access rights in one variable** → saves 1 byte of space, **Block bitmap is one block** → Limits the size of a block group, thus improvising on purpose of a group, **Inode bitmap is one block** → limits total number of files that can belong to a group, File Name is padded → aligns all memory accesses on word boundary, improving performance, **rec_len field in directory entry** → allows holes and linking of entries in directory, **Inode table location in Group Descriptor** → Obvious, as it's per group and not per file-system, **Inode table** → All inodes are kept together so that one disk read leads to reading many inodes together, effectively doing a buffering of subsequent inode reads, and to save space on disk, **Used directories count in group descriptor** →

attempt is made to evenly spread the first-level directories, this count is used there, **Free blocks count in superbloc and group descriptor** → Redundancy to help fsck restore consistency, **Many copies of Superblock** → Redundancy to ensure the most crucial data structure is not lost

Question **6**

Partially correct

Mark 0.43 out of 0.50

Given below are statements about concurrency and parallelism

Select T/F

A concurrent system can allow more than one task to progress, whereas a parallel system can perform more than one task at the same time.

| True | False | | |
|----------------------------------|----------------------------------|--|---|
| <input type="radio"/> | <input checked="" type="radio"/> | Parallel systems allow more than one task to progress while concurrent systems do not. | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | It is possible to have concurrency without parallelism | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | A concurrent system can allow more than one task to progress, whereas a parallel system can perform more than one task at the same time. | ✓ |
| <input checked="" type="radio"/> | <input checked="" type="radio"/> | A concurrent system allows more than one task to progress while a parallel system does not. | ✗ |
| <input type="radio"/> | <input checked="" type="radio"/> | Both concurrency and parallelism are the same. | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | It is possible to have parallelism without concurrency | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | It is not possible to have concurrency without parallelism. | ✓ |

Parallel systems allow more than one task to progress while concurrent systems do not.: False

It is possible to have concurrency without parallelism: True

A concurrent system can allow more than one task to progress, whereas a parallel system can perform more than one task at the same time.: True

A concurrent system allows more than one task to progress while a parallel system does not.: False

Both concurrency and parallelism are the same.: False

It is possible to have parallelism without concurrency: False

It is not possible to have concurrency without parallelism.: False

Question **7**

Correct

Mark 0.50 out of 0.50

Which of the following is not a task of the code of swtch() function

- ☐ a. Switch stacks
- ☒ b. Save the return value of the old context code ✓
- ☐ c. Load the new context
- ☐ d. Jump to next context EIP
- ☒ e. Change the kernel stack location ✓
- ☐ f. Save the old context

The correct answers are: Save the return value of the old context code, Change the kernel stack location

Question **8**

Correct

Mark 0.50 out of 0.50

Which of the following is DONE by allocproc() ?

- ☐ a. ensure that the process starts in trapret()
- ☒ b. ensure that the process starts in forkret() ✓
- ☒ c. allocate PID to the process ✓
- ☒ d. Select an UNUSED struct proc for use ✓
- ☐ e. setup the contents of the trapframe of the process properly
- ☒ f. allocate kernel stack for the process ✓
- ☒ g. setup the trapframe and context pointers appropriately ✓
- ☐ h. setup kernel memory mappings for the process

The correct answers are: Select an UNUSED struct proc for use, allocate PID to the process, allocate kernel stack for the process, setup the trapframe and context pointers appropriately, ensure that the process starts in forkret()

Question 9

Correct

Mark 0.50 out of 0.50

Mark the statements as True or False, w.r.t. passing of arguments to system calls in xv6 code.

| True | False | | |
|----------------------------------|----------------------------------|--|---|
| <input checked="" type="radio"/> | <input type="radio"/> | Integer arguments are copied from user memory to kernel memory using <code>argint()</code> | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | The functions like <code>argint()</code> , <code>argstr()</code> make the system call arguments available in the kernel. | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | The arguments to system call are copied to kernel stack in <code>trapasm.S</code> | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | The arguments to system call originally reside on process stack. | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | Integer arguments are stored in <code>eax</code> , <code>ebx</code> , <code>ecx</code> , etc. registers | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | The arguments are accessed in the kernel code using <code>esp</code> on the trapframe. | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | String arguments are first copied to trapframe and then from trapframe to kernel's other variables. | ✓ |

Integer arguments are copied from user memory to kernel memory using `argint()`: True

The functions like `argint()`, `argstr()` make the system call arguments available in the kernel.: True

The arguments to system call are copied to kernel stack in `trapasm.S`: False

The arguments to system call originally reside on process stack.: True

Integer arguments are stored in `eax`, `ebx`, `ecx`, etc. registers: False

The arguments are accessed in the kernel code using `esp` on the trapframe.: True

String arguments are NOT copied in kernel memory, but just pointed to by a kernel memory pointer: True

String arguments are first copied to trapframe and then from trapframe to kernel's other variables.: False

In the code below assume that each function can be executed concurrently by many threads/processes. Ignore syntactical issues, and focus on the semantics.

This program is an example of

```
spinlock a, b; // assume initialized
thread1() {
    spinlock(a);
    //some code;
    spinlock(b);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
thread2() {
    spinlock(a);
    //some code;
    spinlock(b);
    //some code;
    spinunlock(b);
    spinunlock(a);
}
```

- ☐ a. Deadlock or livelock depending on actual race
- ☐ b. Self Deadlock
- ☐ c. Livelock
- ☒ d. Deadlock ❌
- ☐ e. None of these

Your answer is incorrect.

The correct answer is: None of these

Question 11

Correct

Mark 0.50 out of 0.50

Select the correct statements about paging (not demand paging) mechanism

Select one or more:

- ☐ a. Page table is accessed by the OS as part of execution of an instruction
- ☒ b. OS creates the page table for every process ✓
- ☐ c. User process can update it's own PTBR
- ☒ d. Page table is accessed by the MMU as part of execution of an instruction ✓
- ☒ e. An invalid entry on a page means, it was an illegal memory reference ✓
- ☐ f. An invalid entry on a page means, either it was illegal memory reference or the page was not present in memory.
- ☒ g. The PTBR is loaded by the OS ✓
- ☐ h. User process can update it's own page table entries

Your answer is correct.

The correct answers are: OS creates the page table for every process, The PTBR is loaded by the OS, Page table is accessed by the MMU as part of execution of an instruction, An invalid entry on a page means, it was an illegal memory reference

Question 12

Correct

Mark 0.50 out of 0.50

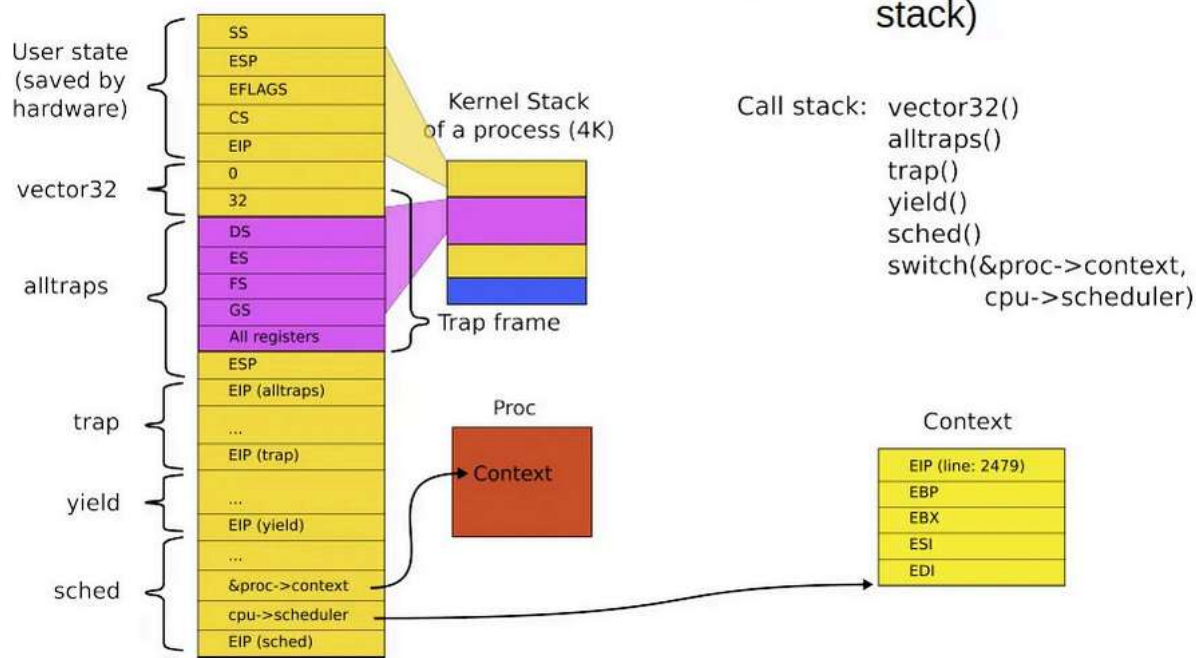
Why V2P_WO is used in entry.S and not V2P ?

- ☐ a. It's a mistake. They could have used the same macro in both places.
- ☐ b. The two macros are different. They lead to different calculations.
- ☐ c. Because the processor can not do a type casting at run time
- ☒ d. Because entry.S is an assembly code file and assemblers do not know about data types and type casting. ✓
- ☐ e. The typecasting has the effect of creating virtual address, while without typecast we get physical address.

Your answer is correct.

The correct answer is: Because entry.S is an assembly code file and assemblers do not know about data types and type casting.

Mark statements as True/False, w.r.t. the given diagram



| True | False | | |
|------------------------------------|------------------------------------|---|---|
| <input type="radio"/> ✗ | <input checked="" type="radio"/> ✓ | The diagram is wrong because it shows the user stack and kernel stack together (continuous), but in practice they are separate | ✓ diagram shows only kernel stack |
| <input checked="" type="radio"/> ✓ | <input type="radio"/> ✗ | The diagram is correct | ✓ |
| <input type="radio"/> ✗ | <input checked="" type="radio"/> ✓ | The "context" yellow coloured box, pointed to by cpu->scheduler is on the kernel stack of the scheduler. | ✓ |
| <input checked="" type="radio"/> ✓ | <input type="radio"/> ✗ | The blue shaded part in "kernel stack of a process(4k)" refers to remaining part of stack (not used yet) | ✓ |
| <input type="radio"/> ✗ | <input checked="" type="radio"/> ✓ | This is a diagram of switch() called from scheduler() | ✓ No. diagram of switch() called from sched() |
| <input checked="" type="radio"/> ✓ | <input type="radio"/> ✗ | The "ESP" (second entry from top) is stack pointer of user-stack of process, while the "ESP" (first entry below pink region) is the trapframe pointer on kernel stack of process. | ✓ |

The diagram is wrong because it shows the user stack and kernel stack together (continuous), but in practice they are separate: False

The diagram is correct: True

The "context" yellow coloured box, pointed to by `cpu->scheduler` is on the kernel stack of the scheduler.: False

The blue shaded part in "kernel stack of a process(4k)" refers to remaining part of stack (not used yet): True

This is a diagram of swtch() called from scheduler(): False

The "ESP" (second entry from top) is stack pointer of user-stack of process, while the "ESP" (first entry below pink region) is the trapframe pointer on kernel stack of process.: True

Question **14**

Correct

Mark 0.50 out of 0.50

Mark the statements as True/False, with respect to the use of the variable "chan" in struct proc.

| True | False | | |
|----------------------------------|----------------------------------|--|---|
| <input type="radio"/> | <input checked="" type="radio"/> | chan is the head pointer to a linked list of processes, waiting for a particular event to occur | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | When chan is not NULL, the 'state' in struct proc must be SLEEPING | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | chan stores the address of the variable, representing a condition, for which the process is waiting. | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | The value of 'chan' is changed only in sleep() | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | in xv6, the address of an appropriate variable is used as a "condition" for a waiting process. | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | Changing the state of a process automatically changes the value of 'chan' | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | 'chan' is used only by the sleep() and wakeup1() functions. | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | when chan is NULL, the 'state' in proc must be RUNNABLE. | ✓ |

chan is the head pointer to a linked list of processes, waiting for a particular event to occur: False

When chan is not NULL, the 'state' in struct proc must be SLEEPING: True

chan stores the address of the variable, representing a condition, for which the process is waiting.: True

The value of 'chan' is changed only in sleep(): True

in xv6, the address of an appropriate variable is used as a "condition" for a waiting process.: True

Changing the state of a process automatically changes the value of 'chan': False

'chan' is used only by the sleep() and wakeup1() functions.: True

when chan is NULL, the 'state' in proc must be RUNNABLE.: False

Question **15**

Correct

Mark 0.50 out of 0.50

Doing a lookup on the pathname /a/b/b/c/d for opening the file "d" requires reading ✓ no. of inodes. Assume that there are no hard/soft links on the path.

Write the answer as a number.

The correct answer is: 6

Question **16**

Partially correct

Mark 0.25 out of 0.50

Suppose a file is to be created in an ext2 file system, in an existing directory /a/b/. Select from below, the list of blocks that may need modification.

Select one or more:

- ☒ a. inode of /a/b/ ✓
- ☐ b. inode of /a/
- ☐ c. link count on /a/b/ inode
- ☐ d. inode bitmap in some block group
- ☒ e. group descriptor(s) ✓
- ☐ f. inode bitmap referring to /a/b/
- ☐ g. new data block in some block group
- ☒ h. superblock ✓
- ☐ i. data blocks of /a/
- ☐ j. inode table in some block group
- ☒ k. block bitmap in some block group ✓
- ☐ l. existing data blocks of /a/b/

Your answer is partially correct.

You have correctly selected 4.

The correct answers are: superblock, group descriptor(s), inode of /a/b/, existing data blocks of /a/b/, inode table in some block group, inode bitmap in some block group, block bitmap in some block group, new data block in some block group

Select the correct statements about hard and soft links

Select one or more:

- ☐ a. Soft link shares the inode of actual file
- ☐ b. Deleting a soft link deletes only the actual file
- ☐ c. Deleting a soft link deletes the link, not the actual file
- ☒ d. Deleting a soft link deletes both the link and the actual file ✖
- ☒ e. Hard links share the inode ✔
- ☒ f. Hard links enforce separation of filename from it's metadata in on-disk data structures. ✔
- ☒ g. Soft links can span across partitions while hard links can't ✔
- ☐ h. Hard links can span across partitions while soft links can't
- ☒ i. Deleting a hard link deletes the file, only if link count was 1 ✔
- ☐ j. Soft links increase the link count of the actual file inode
- ☒ k. Hard links increase the link count of the actual file inode ✔
- ☐ l. Deleting a hard link always deletes the file

Your answer is partially correct.

You have correctly selected 5.

The correct answers are: Soft links can span across partitions while hard links can't, Hard links increase the link count of the actual file inode, Deleting a soft link deletes the link, not the actual file, Deleting a hard link deletes the file, only if link count was 1, Hard links share the inode, Hard links enforce separation of filename from it's metadata in on-disk data structures.

It is proposed that when a process does an illegal memory access, xv6 terminate the process by printing the error message "Illegal Memory Access". Select all the changes that need to be done to xv6 for this as True (Note that the changes proposed here may not cover the exhaustive list of all changes required) and the unnecessary/wrong changes as False.

| Required | Un-necessary/Wrong | | |
|----------------------------------|----------------------------------|---|---|
| <input checked="" type="radio"/> | <input type="radio"/> | Change in the Makefile and instruct cc/ld to start the code of each program at some address other than 0 | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | Change allocuvn() to call mappages() with proper permissions on each page table entry | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | Change mappages() to set specified permissions on each page table entry | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | Mark each page as readonly in the page table mappings | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | Ensure that the address 0 is mapped to invalid | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | Handle the Illegal memory access trap in trap() function, and terminate the currently running process. | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | Change exec to treat text/data sections separately and call allocuvn() with proper flags for page table entries | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | Add code that checks if the illegal memory access trap was due to an actual illegal memory access. | ✓ |

Change in the Makefile and instruct cc/ld to start the code of each program at some address other than 0: Required

Change allocuvn() to call mappages() with proper permissions on each page table entry: Required

Change mappages() to set specified permissions on each page table entry: Un-necessary/Wrong

Mark each page as readonly in the page table mappings: Un-necessary/Wrong

Ensure that the address 0 is mapped to invalid: Required

Handle the Illegal memory access trap in trap() function, and terminate the currently running process.: Required

Change exec to treat text/data sections separately and call allocuvn() with proper flags for page table entries: Required

Add code that checks if the illegal memory access trap was due to an actual illegal memory access.: Un-necessary/Wrong

Question 19

Correct

Mark 0.50 out of 0.50

The variable 'end' used as argument to kinit1 has the value

- ☐ a. 8010a48c
- ☐ b. 80110000
- ☐ c. 80000000
- ☒ d. 801154a8 ✓
- ☐ e. 80102da0
- ☐ f. 81000000

The correct answer is: 801154a8

Question 20

Partially correct

Mark 0.33 out of 0.50

Mark statements about deadlocks as True or false

| True | False | | |
|------------------------------------|------------------------------------|--|---|
| <input type="radio"/> ✗ | <input checked="" type="radio"/> ✓ | Deadlocks are the same as livelocks | ✓ |
| <input checked="" type="radio"/> ✓ | <input type="radio"/> ✗ | A deadlock is possible only if all the 4 conditions of mutual exclusion, cyclic wait, hold and wait, and no preemption are satisfied | ✓ |
| <input checked="" type="radio"/> ✓ | <input type="radio"/> ✗ | A deadlock necessarily requires a cycle in the resource allocation graph | ✓ |
| <input checked="" type="radio"/> ✓ | <input type="radio"/> ✗ | Cycle in the resource allocation graph does not necessarily mean a deadlock | ✓ |
| <input checked="" type="radio"/> ✗ | <input type="radio"/> ✓ | A deadlock must involve at least two processes | ✗ |
| <input type="radio"/> ✓ | <input checked="" type="radio"/> ✗ | Deadlocks are not possible if there is no race | ✗ |

Deadlocks are the same as livelocks: False

A deadlock is possible only if all the 4 conditions of mutual exclusion, cyclic wait, hold and wait, and no preemption are satisfied: True

A deadlock necessarily requires a cycle in the resource allocation graph: True

Cycle in the resource allocation graph does not necessarily mean a deadlock: True

A deadlock must involve at least two processes: False

Deadlocks are not possible if there is no race: True

Question **21**

Correct

Mark 0.50 out of 0.50

Select the statement that most correctly describes what `setupkvm()` does

- ☐ a. creates a 1-level page table for the use by the kernel, as specified in `kmap[]` global array
- ☒ b. creates a 2-level page table setup with virtual->physical mappings specified in the `kmap[]` global array ✓
- ☐ c. creates a 2-level page table setup with virtual->physical mappings specified in the `kmap[]` global array and makes `kpgdir` point to it
- ☐ d. creates a 2-level page table for the use of the kernel, as specified in `gdt_desc`

The correct answer is: creates a 2-level page table setup with virtual->physical mappings specified in the `kmap[]` global array

Question **22**

Correct

Mark 0.50 out of 0.50

Mark statements as T/F

All statements are in the context of preventing deadlocks.

| True | False | | |
|----------------------------------|----------------------------------|--|---|
| <input checked="" type="radio"/> | <input type="radio"/> | A process holding one resources and waiting for just one more resource can also be involved in a deadlock. | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | Hold and wait means a thread/process holding some locks and waiting for acquiring some. | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | If a resource allocation graph contains a cycle then there is a guarantee of a deadlock | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | Circular wait is avoided by enforcing a lock ordering | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | Deadlock is possible if all the conditions are met at the same time: Mutual exclusion, hold and wait, no pre-emption, circular wait. | ✓ |

A process holding one resources and waiting for just one more resource can also be involved in a deadlock.: True

Mutual exclusion is a necessary condition for deadlock because it brings in locks on which deadlock happens: True

The lock ordering to be followed to avoid circular wait is a code in OS that checks for compliance with decided order: False

Hold and wait means a thread/process holding some locks and waiting for acquiring some.: True

If a resource allocation graph contains a cycle then there is a guarantee of a deadlock: False

Circular wait is avoided by enforcing a lock ordering: True

Deadlock is possible if all the conditions are met at the same time: Mutual exclusion, hold and wait, no pre-emption, circular wait.: True

Will this code work for a `spinlock()` operation? The intention here is to call `compare-and-swap()` only if the lock is not held (the if condition checks for the same).

```
void spinlock(int *lock) {  
    {  
        while (true) {  
            if (*lock == 0) {  
                /* lock appears to be available */  
                if (!compare_and_swap(lock, 0, 1))  
                    break  
            }  
        }  
    }  
}
```

- ☐ a. No, because in the case of both processes succeeding in the "if" condition, both may end up acquiring the lock.
- ☒ b. Yes, because no matter in which order the if-check and compare-and-swap run in multiple processes, only one process will succeed in compare-and-swap() and others will keep looping in while-loop. ✓
- ☐ c. Yes, because there is no race to update the lock variable
- ☐ d. No, because this breaks the atomicity requirement of compare-and-test.

Your answer is correct.

The correct answer is: Yes, because no matter in which order the if-check and compare-and-swap run in multiple processes, only one process will succeed in compare-and-swap() and others will keep looping in while-loop.

Question **24**

Partially correct

Mark 0.38 out of 0.50

The kernel ELF file contains these headers

Program Header:

```
LOAD off 0x00001000 vaddr 0x80100000 paddr 0x00100000 align 2**12
      filesz 0x00007aab memsz 0x00007aab flags r-x
LOAD off 0x00009000 vaddr 0x80108000 paddr 0x00108000 align 2**12
      filesz 0x00002516 memsz 0x0000d4a8 flags rw-
STACK off 0x00000000 vaddr 0x00000000 paddr 0x00000000 align 2**4
      filesz 0x00000000 memsz 0x00000000 flags rwx
```

mark the statemetns as True/False

| True | False | | |
|----------------------------------|----------------------------------|--|---|
| <input checked="" type="radio"/> | <input type="radio"/> | Second header is for Data/Globals | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | First header is for the code/text | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | in bootmain() the third header leads to allocation of no-memory. | ✗ |
| <input checked="" type="radio"/> | <input type="radio"/> | Third header is for stack | ✓ |

Second header is for Data/Globals: True

First header is for the code/text: True

in bootmain() the third header leads to allocation of no-memory.: True

Third header is for stack: True

Consider this program.
Some statements are identified using the // comment at the end.
Assume that = is an atomic operation.

```
#include <stdio.h>
#include <pthread.h>
long c = 0, c1 = 0, c2 = 0, run = 1;
void *thread1(void *arg) {
    while(run == 1) { //E
        c = 10; //A
        c1 = c2 + 5; //B
    }
}
void *thread2(void *arg) {
    while(run == 1) { //F
        c = 20; //C
        c2 = c1 + 3; //D
    }
}
int main() {
    pthread_t th1, th2;
    pthread_create(&th1, NULL, thread1, NULL);
    pthread_create(&th2, NULL, thread2, NULL);
    sleep(2);
    run = 0;
    fprintf(stdout, "c = %ld c1+c2 = %ld c1 = %ld c2 = %ld \n", c, c1+c2, c1, c2);
    fflush(stdout);
}
```

Which statements are part of the critical Section?

| YesNo | | | |
|----------------------------------|----------------------------------|---|---|
| <input type="radio"/> | <input checked="" type="radio"/> | C | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | D | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | A | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | F | ✓ |
| <input checked="" type="radio"/> | <input type="radio"/> | B | ✓ |
| <input type="radio"/> | <input checked="" type="radio"/> | E | ✓ |

- C: No
- D: Yes
- A: No
- F: No
- B: Yes
- E: No

Question **26**

Partially correct

Mark 0.38 out of 0.50

Match the code with it's functionality

S = 0

P1:

Statement1;

Signal(S)

Execution order P1, then P2



P2:

Wait(S)

Statment2;

S1 = 0; S2 = 0;

P2:

Statement1;

Signal(S2);

P1:

Wait(S2);

Statemetn2;

Signal(S1);

Execution order P3, P2, P1



P3:

Wait(S1);

Statement S3;

S = 1

Wait(S)

Critical Section

Signal(S);

Binary Semaphore for mutual exclusion



S = 5

Wait(S)

Critical Section

Signal(S)

Counting semaphore



Your answer is partially correct.

You have correctly selected 3.

The correct answer is: **S = 0**

P1:

Statement1;

Signal(S)

P2:

Wait(S)

Statment2; → Execution order P1, then P2, **S1 = 0; S2 = 0;**

P2:

Statement1;

Signal(S2);

P1:

Wait(S2);

Statemetn2;

Signal(S1);

P3:

Wait(S1);

Statement S3; → Execution order P2, P1, P3, **S = 1**

Wait(S)

Critical Section

Signal(S); → Binary Semaphore for mutual exclusion, **S = 5**

Wait(S)

Question **27**

Correct

Mark 0.50 out of 0.50

In an ext2 file system, if the block size is 4KB and partition size is 128 GB, then the number of block groups will be:

Answer: ✓

size * 1024 * 1024 / 4 --> no of blocks

each group = 8 * 4 * 1024 blocks = 32768 blocks

so size * 1024 * 1024 / (4 * 32768) number of groups

The correct answer is: 1024.00

Question **28**

Correct

Mark 0.50 out of 0.50

The "push 0" in vectors.S is

- ☒ a. Place for the error number value ✓
- ☐ b. To be filled in as the return value of the system call
- ☐ c. A placeholder to match the size of struct trapframe
- ☐ d. To indicate that it's a system call and not a hardware interrupt

The correct answer is: Place for the error number value

Question **29**

Partially correct

Mark 0.25 out of 0.50

Match pairs

| | | |
|-----------|---|---|
| mutex | <input type="text" value="atomic test and set with loop"/> | ✗ |
| peterson | <input type="text" value="per process flag, global turn variable"/> | ✓ |
| semaphore | <input type="text" value="wait() and signal()"/> | ✓ |
| spinlock | <input type="text" value="lock() and unlock()"/> | ✗ |

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: mutex → lock() and unlock(), peterson → per process flag, global turn variable, semaphore → wait() and signal(), spinlock → atomic test and set with loop

Question **30**

Correct

Mark 0.50 out of 0.50

Which of the following is done by mappages()?

- ☐ a. allocate page frame if required
- ☐ b. allocate page directory if required
- ☒ c. create page table mappings to the range given by "pa" and "pa + size" ✓
- ☒ d. allocate page table if required ✓
- ☒ e. create page table mappings for the range given by "va" and "va + size" ✓

The correct answers are: create page table mappings for the range given by "va" and "va + size", allocate page table if required, create page table mappings to the range given by "pa" and "pa + size"

◀ Quiz-1 (15 Marks)

Jump to...

ESE(60 Marks) ▶